

Unsupervised and Probabilistic Learning

Coursework

October 2024

1 Models for binary vectors

a) In the question $x^{(n)}$ is a D dimensional random vector that represents a binary pixel value of an image. Hence, for example, a square image the dimensions are $\sqrt{D} \times \sqrt{D}$. A multivariate Gaussian would follow:

$$p(\mathbf{x}) \propto \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \Sigma^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$

where: $\boldsymbol{\mu}$ is the mean vector of length D , Σ is the $D \times D$ covariance matrix, This is inappropriate for the dataset as it has a different sample space that takes binary values of 0 or 1 and, therefore, is bound in this range which does not match the \mathbb{R}^D state space of the Gaussian. This means the Gaussian would assign non-zero probabilities to extreme values that are well beyond this range and are not achievable. The Gaussian model also assumes a continuous distribution that is maximised at the mean. The Gaussian model would therefore not be able to capture the binary distribution of data and would give high weighting to values around 0.5 that are not possible. The shapes are also fairly different as a Gaussian would be unimodal. Lastly, for binary variables, the covariance is bounded and depends on the joint probabilities of pixel occurrences, limiting the covariance matrix to a constrained range that the Gaussian distribution cannot adequately represent.

b) We now model the images as i.i.d from a multivariate Bernoulli distribution with parameter $p = (p_1, \dots, p_D)$:

the likelihood function for the dataset $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$ is:

$$L = \prod_{n=1}^N P(\mathbf{x}^{(n)}|\mathbf{p}) = \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1-x_d^{(n)}}$$

Taking the log of the likelihood we obtain:

$$\log L = \sum_{n=1}^N \sum_{d=1}^D \left[x_d^{(n)} \log p_d + (1 - x_d^{(n)}) \log(1 - p_d) \right]$$

To find the ML estimates, we would maximise the likelihood wrt. p_d however, as log is increasing with positive x values, the maximum stays the same. Hence, we take the partial derivative of the log-likelihood wrt. p_d and set this to 0:

$$\frac{\partial \log L}{\partial p_d} = \sum_{n=1}^N \left(\frac{x_d^{(n)}}{p_d} - \frac{1 - x_d^{(n)}}{1 - p_d} \right) = 0$$

Note that the partial wrt. a given p_d goes to 0 for all other ds hence the sum over d disappears. Solving:

$$\begin{aligned} \sum_{n=1}^N x_d^{(n)}(1 - p_d) &= \sum_{n=1}^N (1 - x_d^{(n)})p_d \\ \sum_{n=1}^N x_d^{(n)} - p_d \sum_{n=1}^N x_d^{(n)} &= p_d \sum_{n=1}^N (1 - x_d^{(n)}) \\ \sum_{n=1}^N x_d^{(n)} &= p_d \left(\sum_{n=1}^N x_d^{(n)} + \sum_{n=1}^N (1 - x_d^{(n)}) \right) = p_d N \end{aligned}$$

So the ML estimate for p_d is:

$$\implies \hat{p}_d^{ML} = \frac{1}{N} \sum_{n=1}^N x_d^{(n)}$$

We can interpret this maximum likelihood estimate of p_d as the average value of the d -th pixel across all N images.

c) Assuming Beta priors on p_d :

$$P(p_d) = \frac{1}{B(\alpha, \beta)} p_d^{\alpha-1} (1 - p_d)^{\beta-1}$$

for $P(p) = \prod_d P(p_d)$. To find the MAP estimator want to, effectively, find the most probable parameter given the data. \hat{p} = the maximisation of $P(p|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)})$.

The posterior distribution is proportional to the product of the likelihood and the prior under Bayes' Theorem:

$$\begin{aligned} P(p|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}) &\propto L \times P(\mathbf{p}) = \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1 - x_d^{(n)}} \cdot \prod_{d=1}^D p_d^{\alpha-1} (1 - p_d)^{\beta-1} \\ &= \prod_{d=1}^D p_d^{\sum_{n=1}^N x_d^{(n)} + \alpha - 1} (1 - p_d)^{N - \sum_{n=1}^N x_d^{(n)} + \beta - 1} \end{aligned}$$

The denominator is independent of \mathbf{p} so that term can be ignored as it will go to 0 upon maximising. To find the MAP estimator, we maximise the posterior

wrt. \mathbf{p} . Since the priors are independent across dimensions, the MAP estimate for each p_d can be found separately.

Using a similar technique as in part **b**:

$$\log P(\mathbf{p}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}) = \sum_{d=1}^D \left(\left(\sum_{n=1}^N x_d^{(n)} + \alpha - 1 \right) \log p_d + \left(N - \sum_{n=1}^N x_d^{(n)} + \beta - 1 \right) \log(1 - p_d) \right)$$

Taking the derivative with respect to p_d and setting to zero:

$$\frac{\partial}{\partial p_d} \log P(\mathbf{p}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}) = \frac{\sum_{n=1}^N x_d^{(n)} + \alpha - 1}{p_d} - \frac{N - \sum_{n=1}^N x_d^{(n)} + \beta - 1}{1 - p_d} = 0$$

Solving and rearranging:

$$\begin{aligned} \frac{\sum_{n=1}^N x_d^{(n)} + \alpha - 1}{p_d} &= \frac{N - \sum_{n=1}^N x_d^{(n)} + \beta - 1}{1 - p_d} \\ \left(\sum_{n=1}^N x_d^{(n)} + \alpha - 1 \right) (1 - p_d) &= \left(N - \sum_{n=1}^N x_d^{(n)} + \beta - 1 \right) p_d \\ \sum_{n=1}^N x_d^{(n)} + \alpha - 1 - p_d \left(\sum_{n=1}^N x_d^{(n)} + \alpha - 1 \right) &= p_d \left(N - \sum_{n=1}^N x_d^{(n)} + \beta - 1 \right) \\ \sum_{n=1}^N x_d^{(n)} + \alpha - 1 &= p_d (N + \alpha + \beta - 2) \\ \implies \hat{p}_d^{\text{MAP}} &= \frac{\sum_{n=1}^N x_d^{(n)} + \alpha - 1}{N + \alpha + \beta - 2} \end{aligned}$$

d) and e) please find the code below for parts d and e (including the given functions for displaying the data)

```

1 import numpy as np
2 from matplotlib import pyplot as plt
3
4 def ml_parameters(Y):
5     """
6     Compute the ML estimates for the Bernoulli parameters.
7     parameters:
8         Y
9     returns:
10         params_ml: Array of ML estimates for each pixel
11     """
12     N, D = Y.shape
13     S_d = np.sum(Y, axis=0)
14     params_ml = S_d / N
15 
```

```

16     return params_ml
17
18 def map_parameters(Y, alpha=3, beta=3):
19     """
20     Compute the MAP estimates for the Bernoulli parameters
21     with Beta priors.
22     parameters:
23         Y: N x D binary matrix.
24         alpha: beta distribution param
25         beta: beta distribution param
26     returns:
27         params_map: Array of MAP estimates
28     """
29     N, D = Y.shape
30     S_d = np.sum(Y, axis=0) #sum of each column
31     params_map = (S_d + alpha - 1) / (N + alpha + beta - 2)
32     return params_map
33
34 def display_image(params, title):
35     """
36     Display a D-dimensional parameter vector as an 8x8 image
37     with a colour scale
38     parameters:
39         params: Parameter vector
40         title: image title
41     """
42     plt.figure(figsize=(4,4))
43     img=np.reshape(params, (8,8))
44     plt.imshow(img, interpolation="none", cmap='winter')
45     plt.title(title)
46     plt.colorbar()
47     plt.show()
48
49 def plot_difference(params_ml, params_map):
50     """
51     Plot the difference between ML and MAP parameters
52     parameters:
53         params_ml: ML estimates
54         params_map: MAP estimates
55     """
56     difference = params_ml - params_map
57     plt.figure(figsize=(4,4))
58     plt.imshow(np.reshape(difference, (8,8)), interpolation=
59         "none", cmap='bwr')
60     plt.title('Difference')
61
62     plt.colorbar()
63     plt.show()
64
65 def main():

```

```

63 Y = load_data('/Users/baidn/Downloads/binarydigits.txt')
64 N, D = Y.shape
65
66 #compute ml and map
67 params_ml = ml_parameters(Y)
68 alpha, beta_val = 3, 3
69 params_map = map_parameters(Y, alpha, beta_val)
70
71 # plot the images and difference
72 display_image(params_ml, title='ML Parameters')
73 display_image(params_map, title='MAP Parameters (alpha
74                 =3, beta=3)')
75
76 plot_difference(params_ml, params_map)
77
78 if __name__ == "__main__":
79     main()

```

Listing 1: Python Code for Learning ML and MAP Parameters

Visualisation of Learned Parameters




Figure 1: Learned ML Parameters

The MAP estimator integrates prior beliefs about the parameters with the observed data, balancing them based on the chosen prior parameters. Specifically, with $\alpha = \beta = 3$, the prior suggests a belief that each pixel has an inherent probability of 0.5 of being active, providing a regularisation effect that can prevent extreme estimates in scenarios with sparse or noisy data. For example, if a pixel was black in all the images the ML would be 1, hence the prior allows for some noise as you will not have a probability that is exactly 1 or 0.

As the size of the data increases the MAP and ML become increasingly similar as the prior becomes less significant. The Difference is very small as




Figure 2: Learned MAP Parameters




Figure 3: Difference: ML-MAP

shown in Fig 3.

The introduction of prior information can bias the estimates away from the true ML estimates, particularly if the prior is not well-aligned with the underlying data distribution so in this case a ML estimate may be better.

In our case our MAP parameters are biased to be closer to 0.5 which is shown by the difference plot in Fig 3. as the difference is positive for ML greater than 0.5 and negative for below.

2 Model selection

We aim to find the normalised posterior probability for each model given the data. This is given by

$$P(p^a \mid \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}) = \frac{P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \mid p^a)P(p)}{P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)})}$$

We are told that each distribution has equal priors $P(p)$ (i.e.=1/3) and $P(\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\})$ is the same for each model hence for the relative/normalised probabilities we only need to find $P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \mid p^a) \mid p_d = 0.5$

In this model, each component x_d of the data is generated from a Bernoulli distribution with $p_d = 0.5$. Thus, the probability of observing a specific dataset $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$ with N binary vectors, each with D components, is given by

$$P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \mid p^a) = \prod_{n=1}^N \prod_{d=1}^D P(x_d^{(n)} \mid p_d = 0.5) = \prod_{n=1}^N \prod_{d=1}^D (0.5)^{x_d^{(n)}} (0.5)^{1-x_d^{(n)}} = \prod_{n=1}^N \prod_{d=1}^D 0.5$$

Since there are $N \times D$ entries:

$$P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \mid p^a) = 0.5^{N \times D}.$$

Taking the log

$$\log P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \mid p^a) = (N \times D) \cdot \log(0.5).$$

b) Identical with unknown p_d

Here we assume that all components share the same unknown probability $p_d = p$, which is inferred from the data.

Let S be the total count of 1s and Y be the total count of 0s across all data):

$$S = \sum_{n=1}^N \sum_{d=1}^D x_d^{(n)}, Y = (N \times D) - S.$$

The likelihood for S successes and Y failures given p is $p^S(1-p)^Y$. The posterior probability with a uniform prior is given by the Beta function from the following integral:

$$P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \mid p^b) = \int_0^1 p^S (1-p)^Y dp = \frac{\Gamma(S+1)\Gamma(Y+1)}{\Gamma(S+Y+2)}.$$

Taking the log-probability, we have

$$\log P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \mid p^b) = \log \left(\frac{\Gamma(S+1)\Gamma(Y+1)}{\Gamma(S+Y+2)} \right) = \text{betaln}(S+1, Y+1).$$

c) In this model, each component x_d has its own independent Bernoulli distribution with separate unknown probability p_d , which follows a uniform prior distribution. We assume a Beta prior for each p_d .

For each component d , let S_d represent the total number of 1s observed in the d -th column of the data set and $Y_d = N - S_d$ be the number of 0s. The posterior probability for each p_d is therefore

$$P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} | p^c) = \prod_{d=1}^D \int_0^1 p_d^{S_d} (1 - p_d)^{Y_d} dp_d = \prod_{d=1}^D \frac{\Gamma(S_d + 1) \Gamma(Y_d + 1)}{\Gamma(N + 2)}.$$

Taking the log-probability for the whole dataset, we get

$$\log P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} | p^c) = \sum_{d=1}^D \text{betaln}(S_d + 1, Y_d + 1).$$

To determine the relative posterior probabilities for each model, we combine the log-probabilities calculated above. Assuming equal prior probabilities for each model, the posterior probabilities are proportional to the likelihoods, so we normalise by summing the exponentiated log-probabilities across all models:

$$P(p^i | \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}) = \frac{\exp(\log P(\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\} | p_i))}{\sum_{j=a,b,c} \exp(\log P(\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\} | p_j))}.$$

Due to small probabilities we use logs and after trying multiple methods we use the `logsumexp` python function to find these absolute probabilities.

```

1  Y = np.loadtxt('/Users/baidn/Downloads/binarydigits.txt')
2  )
3  N, D = Y.shape
4
5  # Model a
6  m1 = np.log(0.5) * (N * D)
7
8  # Model b
9  S = np.sum(Y) # Total 1s
10 Y = N * D - S # Total 0s
11 m2 = beta_func(S + 1, Y + 1) # Log probability with
    Beta prior
12
13 # Model c
14 S_d = np.sum(Y, axis=0) # Count of 1s for each
    component
15 m3 = np.sum(beta_func(S_d + 1, N - S_d + 1)) # Sum log
    probabilities for each p_d
16
17 # Normalise
18 log_probs = [m1, m2, m3]
19 total_log_prob = logsumexp(log_probs)
20
21 print("relative posterior log relative and absolute
    probabilities :")
22 print(f"Model a: {m1 - total_log_prob} and {np.exp(m1 -
    total_log_prob)}")

```



```
relative posterior log and absolute probabilities:
Model a: -584.9462116625186 and 9.142986218361563e-255
Model b: -432.52559865622743 and 1.4339011785434019e-188
Model c: 0.0 and 1.0 -1.4339011785434019e-188 -9.142986218361563e-255
```

Figure 4: relative log and absolute probabilities

```
23 | print(f"Model b: {m2 - total_log_prob} and {np.exp(m2 -
    | total_log_prob)}")
24 | print(f"Model c: {m3 - total_log_prob} and {np.exp(m3 -
    | total_log_prob)}")
```

Listing 2: Model Selection with Bernoulli Distributions

3 EM for Binary Data

a) The likelihood for a mixture model consisting of K multivariate Bernoulli distributions where π_i represents the mixing proportions such that $\sum_i \pi_i = 1$ can be found by considering an image $\mathbf{x}^{(n)}$.

Each image $\mathbf{x}^{(n)} \in \{0, 1\}^D$ (for $n = 1, \dots, N$) is i.i.d under the mixture model. The probability of an image $\mathbf{x}^{(n)}$ given component k is defined as a product of Bernoulli distributions for each pixel, assuming pixel independence within each component.

$$P(\mathbf{x}^{(n)} \mid \pi, \mathbf{P}) = \sum_k \pi_k P(\mathbf{x}^{(n)} \mid \mathbf{p}_k) = \sum_{k=1}^K \pi_k \prod_{d=1}^D p_{kd}^{x_d^{(n)}} (1 - p_{kd})^{1-x_d^{(n)}},$$

where \mathbf{P} matrix of Bernoulli parameters with elements where p_{kd} is the probability that pixel d takes the value 1 under component k and \mathbf{p}_k is the k th row in the \mathbf{P} matrix.

The likelihood of the entire dataset is given by

$$P(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)} \mid \pi, P) = \prod_{n=1}^N \sum_{k=1}^K \pi_k \prod_{d=1}^D p_{kd}^{x_d^{(n)}} (1 - p_{kd})^{1-x_d^{(n)}}.$$

Taking the log:

$$\log P(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)} \mid \pi, P) = \sum_{n=1}^N \log \left(\sum_{k=1}^K \pi_k \prod_{d=1}^D p_{kd}^{x_d^{(n)}} (1 - p_{kd})^{1-x_d^{(n)}} \right).$$

b) The E-step involves computing the responsibility of each mixture component k for a given vector $\mathbf{x}^{(n)}$. Using Bayes' theorem

$$r_{nk} = P(s^n = k \mid \mathbf{x}^{(n)}, \pi, \mathbf{P}) = \frac{P(\mathbf{x}^{(n)} \mid s^{(n)} = k, \mathbf{P}) P(s^{(n)} = k \mid \pi)}{P(\mathbf{x}^{(n)} \mid \pi, \mathbf{P})}$$

where this mixture can be written as a latent variable model with $s^{(n)} \in 1, \dots, K$ and $P(s^{(n)} = k \mid \pi) = \pi_k$

$$\begin{aligned} r_{nk} &= \frac{P(\mathbf{x}^{(n)} \mid s^{(n)} = k, \mathbf{P})P(s^{(n)} = k \mid \pi)}{\sum_{j=1}^K P(s^{(n)} = j \mid \pi)P(\mathbf{x}^{(n)} \mid s^{(n)} = j, \mathbf{P})} \\ &= \frac{\pi_k \prod_{d=1}^D p_{kd}^{x_d^{(n)}} (1 - p_{kd})^{1-x_d^{(n)}}}{\sum_{j=1}^K \pi_j \prod_{d=1}^D p_{jd}^{x_d^{(n)}} (1 - p_{jd})^{1-x_d^{(n)}}}. \end{aligned}$$

c) In this part, we aim to determine the parameters π and \mathbf{P} that maximise the expected log-joint probability of the observed data and the latent variables within the EM algorithm framework. The latent variables are denoted as $s^{(n)}$.

$$\arg \max_{\pi, P} \sum_{n=1}^N \langle \log P(\mathbf{x}^{(n)}, \mathbf{s}^{(n)} \mid \pi, P) \rangle = \arg \max_{\pi, P} \sum_{n=1}^N \langle \log P(\mathbf{x}^{(n)} \mid \mathbf{s}^{(n)}, \pi, P) P(s^{(n)} \mid \pi, \mathbf{P}) \rangle$$

this will help obtain the iterative update for the parameters π and P in the M-step of EM.

In the EM algorithm, the M-step involves maximising the expected complete data log-likelihood with respect to the parameters π and P , given the current estimate of the responsibilities r_{nk} computed in the E-step.

Using the above equation, the expected complete log-likelihood is given by:

$$\mathcal{L} = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \left[\log \pi_k + \sum_{d=1}^D \left(x_d^{(n)} \log p_{kd} + (1 - x_d^{(n)}) \log(1 - p_{kd}) \right) \right],$$

To maximise \mathcal{L} with respect to π_k , we set up the Lagrangian with a multiplier λ with the constraint $\sum_{k=1}^K \pi_k = 1$:

$$\mathcal{L} = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \log \pi_k + \lambda \left(\sum_{k=1}^K \pi_k - 1 \right).$$

Taking the partial derivative with respect to π_k and setting it to zero:

$$\frac{\partial \mathcal{L}}{\partial \pi_k} = \sum_{n=1}^N \frac{r_{nk}}{\pi_k} + \lambda = 0.$$

Solving for π_k :

$$\pi_k = -\frac{1}{\lambda} \sum_{n=1}^N r_{nk}.$$

Applying the constraint $\sum_{k=1}^K \pi_k = 1$:

$$\sum_{k=1}^K \pi_k = -\frac{1}{\lambda} \sum_{k=1}^K \sum_{n=1}^N r_{nk} = -\frac{1}{\lambda} \sum_{n=1}^N \sum_{k=1}^K r_{nk} = -\frac{1}{\lambda} N.$$

Since $\sum_{k=1}^K \pi_k = 1$, we have:

$$-\frac{N}{\lambda} = 1 \implies \lambda = -N.$$

The updated mixing proportions are:

$$\implies \pi_k^{new} = \frac{1}{N} \sum_{n=1}^N r_{nk}.$$

To maximise \mathcal{L} with respect to p_{kd} , we take the derivative of \mathcal{L} with respect to p_{kd} , incorporating the constraint $0 \leq p_{kd} \leq 1$.

The relevant part of \mathcal{L} with p_{kd} (for which derivativ would not go to 0) is:

$$\mathcal{L}_{p_{kd}} = \sum_{n=1}^N r_{nk} \left(x_d^{(n)} \log p_{kd} + (1 - x_d^{(n)}) \log(1 - p_{kd}) \right).$$

Taking the derivative with respect to p_{kd} and setting it to zero:

$$\frac{\partial \mathcal{L}_{p_{kd}}}{\partial p_{kd}} = \sum_{n=1}^N r_{nk} \left(\frac{x_d^{(n)}}{p_{kd}} - \frac{1 - x_d^{(n)}}{1 - p_{kd}} \right) = 0.$$

Simplify the derivative:

$$\sum_{n=1}^N r_{nk} \left(\frac{x_d^{(n)}(1 - p_{kd}) - (1 - x_d^{(n)})p_{kd}}{p_{kd}(1 - p_{kd})} \right) = 0.$$

Multiply both sides by $p_{kd}(1 - p_{kd})$:

$$\sum_{n=1}^N r_{nk} \left(x_d^{(n)}(1 - p_{kd}) - (1 - x_d^{(n)})p_{kd} \right) = 0.$$

Simplifying:

$$\sum_{n=1}^N r_{nk} \left(x_d^{(n)} - x_d^{(n)} p_{kd} - p_{kd} + x_d^{(n)} p_{kd} \right) = 0.$$

$$\sum_{n=1}^N r_{nk} \left(x_d^{(n)} - p_{kd} \right) = 0.$$

Therefore:

$$\sum_{n=1}^N r_{nk} x_d^{(n)} = p_{kd} \sum_{n=1}^N r_{nk}.$$

$$\Rightarrow p_{kd}^{new} = \frac{\sum_{n=1}^N r_{nk} x_d^{(n)}}{\sum_{n=1}^N r_{nk}}.$$

d) the code implements the EM algorithm using the above equations using logs due to small likelihoods. When running the algorithm multiple times the convergence for different Ks seemed to vary quite a lot but in general higher Ks took longer to converge to the threshold. d) Please find the code for the EM algorithm below. I used log-likelihoods and initialised π uniformly and \mathbf{P} was initialised between 0.1 and 0.9 to avoid extreme probabilities to start with.

```

1 import numpy as np
2 from scipy.special import logsumexp
3 from matplotlib import pyplot as plt
4
5 # load data
6 X = np.loadtxt('/Users/baidn/Downloads/binarydigits.txt')
7 N, D = X.shape # N = number of data points, D =
8               # dimensionality
9 K_val = [2, 3, 4, 7, 10] # different Ks
10
11 def EM(K, X, max_iter):
12     """
13     Performs the EM algorithm
14
15     parameters:
16     K: Number of mixture components
17     X: Data matrix of shape (N, D)
18     max_iter: Maximum number of iterations
19
20     returns:
21     pi: Learned mixing proportions of shape (K,)
22     P: Learned Bernoulli parameters of shape (K, D)
23     log_likelihoods: List of log-likelihood values at each
24                     iteration
25     """
26     thrs = 1e-6 # threshold for convergence
27
28     # uniform mixing
29     pi = np.full(K, 1.0 / K)
30
31     # Bernoulli parameters random between 0.1 and 0.9 -
32     # avoid log 1
33     P = np.random.rand(K, D) * 0.8 + 0.1
34
35     def E_step(X, P, pi):
36         """
37         E-step - responsibility
38
39         returns:
40         R: Responsibility

```

```

38     """
39     #log probabilities to prevent underflow
40     log_P = np.log(P)
41     log_1_minus_P = np.log(1 - P)
42
43     # log-likelihood for each component and data point
44     log_likelihood = np.zeros((N, K))
45     for k in range(K):
46         # log-likelihood for component k
47         log_prob = X @ log_P[k] + (1 - X) @
48             log_1_minus_P[k] #
49         log_likelihood[:, k] = log_prob + np.log(pi[k])
50
51     #responsibilities in log space to prevent underflow
52     log_sum = logsumexp(log_likelihood, axis=1, keepdims
53         =True) #
54     log_R = log_likelihood - log_sum
55     R = np.exp(log_R) #convert back
56
57     return R
58
59 def M_step(X, R):
60     """
61     M - Update params pi and P.
62
63     returns:
64     pi_new: Updated mixing proportions
65     P_new: Updated Bernoulli params
66     """
67     epsilon = 1e-7 #prevetn div by 0
68
69     # Update mixing proportions as in c)
70     N_k = np.sum(R, axis=0)
71     pi_new = N_k / N
72
73     # Update Bernoulli parameters as in c)
74     P_new = np.zeros((K, D))
75     for k in range(K):
76         # weighted sum of data points for each feature
77         numerator = R[:, k] @ X
78         denominator = N_k[k] + epsilon # no div by 0
79         P_new[k] = numerator / denominator
80
81     # Ensure P_new values are within (0, 1)
82     P_new = np.clip(P_new, epsilon, 1 - epsilon)
83
84     return pi_new, P_new
85
86 def compute_log_likelihood(X, pi, P):
87     """

```

```

86         computes the log-likelihood
87
88     params:
89         uses current parameters for pi and P
90
91     returns:
92         total_log_likelihood: Scalar value of the total log-
93             likelihood
94         """
95         # log probabilities
96         log_P = np.log(P)
97         log_1_minus_P = np.log(1 - P)
98
99         #log-likelihood for components and data points
100         log_likelihood = np.zeros((N, K))
101         for k in range(K):
102             log_prob = X @ log_P[k] + (1 - X) @
103                 log_1_minus_P[k]
104             log_likelihood[:, k] = log_prob + np.log(pi[k])
105
106         #total log-likelihood using logsumexp
107         log_sum = logsumexp(log_likelihood, axis=1)
108         total_log_likelihood = np.sum(log_sum)
109
110         return total_log_likelihood
111
112 log_likelihoods = []
113
114 for iteration in range(max_iter):
115     #E
116     R = E_step(X, P, pi)
117
118     #M
119     pi, P = M_step(X, R)
120
121     #log-likelihood
122     ll = compute_log_likelihood(X, pi, P)
123     log_likelihoods.append(ll)
124
125     # convergence check
126     if iteration > 0:
127         ll_change = np.abs(log_likelihoods[-1] -
128             log_likelihoods[-2])
129         if ll_change < thrs:
130             print(f"Converged at iteration {iteration}."
131                 )
132             break
133
134     print(f"Iteration {iteration + 1}: Log-Likelihood =
135         {ll:.4f}")

```

```

131         return pi, P, log_likelihoods
132
133
134     # Dictionaries to store results for different K values
135     log_likelihood_histories = {}
136     learned_parameters = {}
137
138     #loop over each K value
139     for K in K_val:
140         print(f"Run {K}")
141
142
143         pi_learned, P_learned, ll_history = EM(K, X, max_iter
144             =100)
145
146         # store history
147         log_likelihood_histories[K] = ll_history
148         learned_parameters[K] = {
149             'pi': pi_learned,
150             'P': P_learned
151         }
152
153         #display mixing
154         print(f"pi for K={K}")
155         print(pi_learned)
156
157     # plot log-likelihoods for each K
158     plt.figure(figsize=(10, 6))
159     for K in K_val:
160         ll_history = log_likelihood_histories[K]
161         iterations = range(1, len(ll_history) + 1)
162         plt.plot(iterations, ll_history, label=f'K = {K}')
163
164     plt.title('Log-likelihood vs iteration for Different K')
165     plt.xlabel('Iteration')
166     plt.ylabel('Log-likelihood')
167     plt.legend()
168     plt.show()
169
170     # plot the learned Bernoulli parameters P for each K
171     for K in K_val:
172         P_learned = learned_parameters[K]['P']
173         num_components = P_learned.shape[0]
174
175         #subplot
176         cols = min(num_components, 5)
177         rows = (num_components + cols - 1) // cols
178         #looping through each if tge mixings to plot the
179         paramter distribution

```

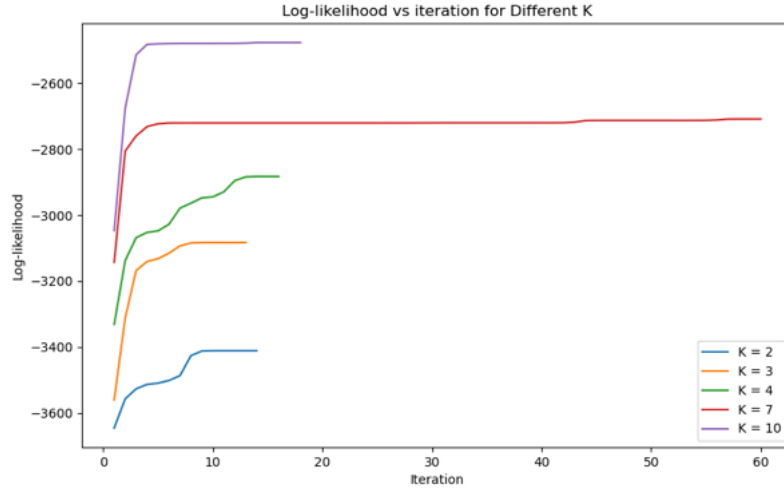


Figure 5: Log-likelihood vs Iteration for varying K

```

179 plt.figure(figsize=(cols * 2, rows * 2))
180 for k in range(num_components):
181     plt.subplot(rows, cols, k + 1)
182     plt.imshow(P_learned[k].reshape(8, 8), interpolation
183               = 'none', cmap='bwr')
184     plt.title(f"K={K}, p = {round(learned_parameters[K
185                               ]['pi'][k], 2)}")
186
187 plt.tight_layout()
188 plt.show()

```

e) I ran the EM a few times with random initialisation. For smaller values of K it seems there is somewhat similar solutions with the parameters looking like 0,5 for 2 of the trials and 0,7 for the other for K=2. Even when running for more than 3 trials the 0 always seems to be one of the parameter images for K=2. These common features become less true as the value of K increases with a considerable amount for variance for K=10 among the trials. We can also observe that for certain digits there is different representations as the zero seems to have a very slanted structure or a centred one. Also, as K increases the EM algorithm is more susceptible to variability across runs due to multiple local optima.

However, a very small K would not allow us to train multivariate distributions for each number as we have 0,5,7 and various variants of these. Hence, maybe with more trials we could find a value of K that is greater than this but less than 10. Hence, we could improve the model by choosing K according to




Figure 6: K=2, trial 1




Figure 7: K=2, trial 2

the number of different variants of numbers in the dataset, probably a K just higher than 3.

We could consider a larger/more diverse dataset to improve our model and make it more flexible to new numbers.

Also, for high values of K it is interesting to see that there are some mixtures that have very low probabilities and their Parameter distribution is relatively binary implying they may only represent a single image. Hence, the model could also be improved by adding intelligent priors to allow for a more representative distribution.

Parameter plots for different values of K:

f) The log likelihoods are converted to bits by $L_b = \frac{\log L}{\ln(2)}$ and the number of bits from the naive encoding is $N \times D$. We can find the average code length per bit by dividing the compressed size and L_b by the encoding. I am unsure




Figure 8: K=2, trial 3




Figure 9: K=3, trial 1




Figure 10: K=3, trial 2




Figure 11: K=3, trial 3




Figure 12: K=4, trial 1




Figure 13: K=4, trial 2




Figure 14: K=4, trial 3




Figure 15: K=7, trial 1




Figure 16: K=7, trial 2




Figure 17: K=7, trial 3




Figure 18: K=10, trial 1




Figure 19: K=10, trial 2




Figure 20: K=10, trial 3

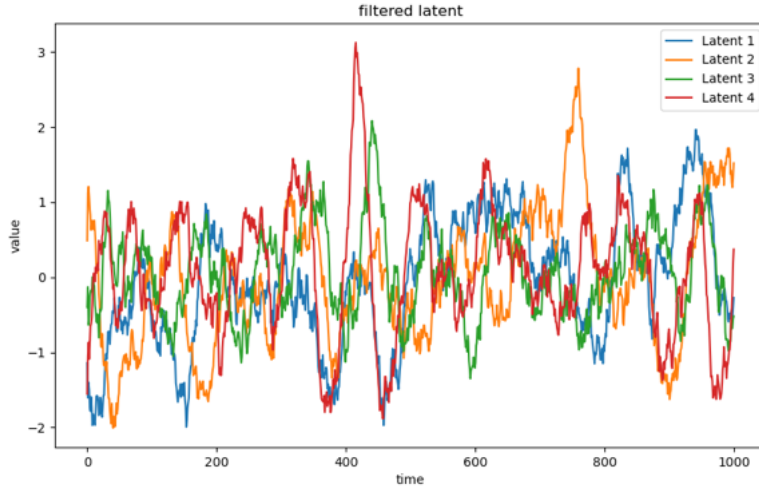


Figure 21: Filtered plot

as to how to find the exact gzip encoding but the average code length per bit increases with K as seen from the log-likelihood graph.

4 LGSSMs, EM and SSID

a) To run the functions on the provided training data, I first had to transpose the loaded matrices. The code defines the parameters as in the questions and plots the data. $k=4$ from matrix A. Following the pseudocode in the question I find the following plots:

b) We try to estimate the parameters \mathbf{A} , \mathbf{Q} , \mathbf{C} and \mathbf{R} given the data we have. Note:

$$\mathbf{y}_t \sim N(\mathbf{A}\mathbf{y}_{t-1}, \mathbf{Q})$$

$$\mathbf{x}_t \sim N(\mathbf{C}\mathbf{y}_t, \mathbf{R})$$




Figure 22: Filtered log determinant

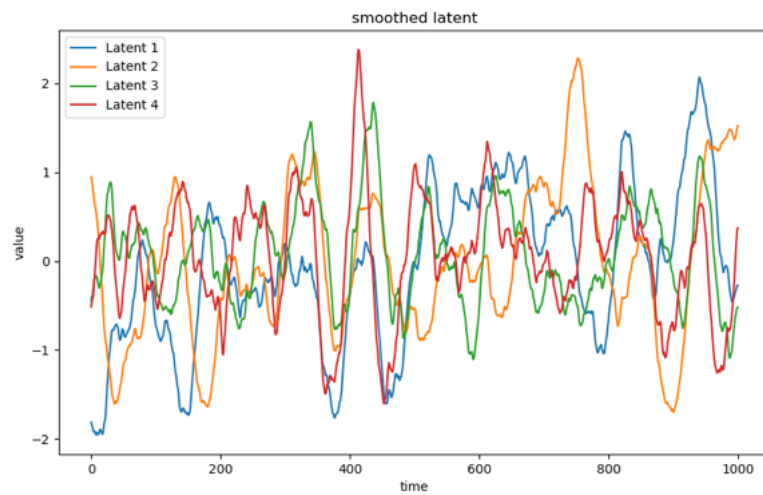


Figure 23: Smoothed plot

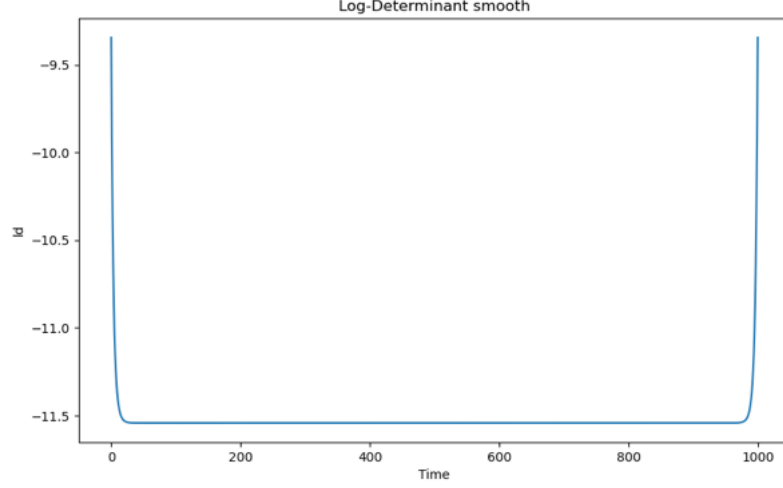


Figure 24: Smoothed log determinant

Showing the M-Step for R:

$$p(x_t|y_t) \propto \exp(-1/2(x_t - Cy_t)^T R^{-1}(x_t - Cy_t))$$

given this the log-likelihood

$$L = \left\langle \sum_{t=1}^T \ln p(x_t|y_t) \right\rangle = -\frac{1}{2} \sum_{t=1}^T \langle (x_t - Cy_t)^T R^{-1}(x_t - Cy_t) \rangle - \frac{T}{2} \ln |R|$$

$$\langle (x_t - Cy_t)(x_t - Cy_t)^T \rangle = x_t x_t^T - x_t \langle y_t^T \rangle C^T - C \langle y_t \rangle x_t^T + C \langle y_t y_t^T \rangle C^T$$

$$L = -\frac{1}{2} Tr(R^{-1}(x_t x_t^T - x_t \langle y_t^T \rangle C^T - C \langle y_t \rangle x_t^T + C \langle y_t y_t^T \rangle C^T)) - \frac{T}{2} \ln |R| + \text{const}$$

minimising R^{-1} maximises R

$$\frac{\partial L}{\partial R^{-1}} = \frac{1}{2} \sum_t (x_t x_t^T - x_t \langle y_t^T \rangle C^T - C \langle y_t \rangle x_t^T + C \langle y_t y_t^T \rangle C^T) - \frac{T}{2} R = 0$$

$$R = \frac{1}{T} \sum_t (x_t x_t^T - x_t \langle y_t^T \rangle C^T - C \langle y_t \rangle x_t^T + C \langle y_t y_t^T \rangle C^T)^T$$

as $\frac{\partial Tr(AB)}{\partial B} = B^T$ Note that the update for C (from lectures) is,

$$C_{\text{new}} = \sum_t x_t \langle y_t \rangle^T \left(\sum_t \langle y_t y_t^T \rangle \right)^{-1}$$

given

$$C_{\text{new}} \sum_t \langle y_t y_t^T \rangle = \sum_t x_t \langle y_t^T \rangle$$

this gives

$$R_{\text{new}} = \frac{1}{T} \left(\sum_{t=1}^T x_t x_t^T - \sum_{t=1}^T x_t \langle y_t^T \rangle C_{\text{new}}^T \right)$$

For Q,

$$p(y_t | y_{t-1}) \propto \exp \left(-\frac{1}{2} (y_t - A y_{t-1})^T Q^{-1} (y_t - A y_{t-1}) \right)$$

$$L = \left\langle \sum_{t=2}^T \ln p(y_t | y_{t-1}) \right\rangle = -\frac{1}{2} \sum_{t=2}^T \langle (y_t - A y_{t-1})^T Q^{-1} (y_t - A y_{t-1}) \rangle - \frac{T-1}{2} \ln |Q| + \text{const}$$

$$\langle (y_t - A y_{t-1})(y_t - A y_{t-1})^T \rangle = \langle y_t y_t^T \rangle - A \langle y_{t-1} y_t^T \rangle - \langle y_t y_{t-1}^T \rangle A^T + A \langle y_{t-1} y_{t-1}^T \rangle A^T$$

Maximising with Respect to Q (same as before)

$$\frac{\partial L}{\partial Q^{-1}} = \frac{1}{2} (\langle y_t y_t^T \rangle - A \langle y_{t-1} y_t^T \rangle - \langle y_t y_{t-1}^T \rangle A^T + A \langle y_{t-1} y_{t-1}^T \rangle A^T) - \frac{T-1}{2} Q = 0$$

$$Q = \frac{1}{T-1} (\langle y_t y_t^T \rangle - A \langle y_{t-1} y_t^T \rangle - \langle y_t y_{t-1}^T \rangle A^T + A \langle y_{t-1} y_{t-1}^T \rangle A^T)$$

Using update for A in lecture notes and using a similar method as above to cancel terms Final update for Q:

$$Q_{\text{new}} = \frac{1}{T-1} \left(\sum_{t=2}^T \langle y_t y_t^T \rangle - A_{\text{new}} \left(\sum_{t=2}^T \langle y_t y_{t-1}^T \rangle \right)^T \right)$$

The code includes the plotting of the figures in a), the Kalman function given to us and the EM

When run, I obtain a plot that shows the when starting with the generating initialisations there is much faster convergence compared to random - which makes sense as these are generated parameters. Also, the generating starts off with a much higher log-likelihood.

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4
5 X = np.loadtxt('/Users/baidn/Downloads/ssm_spins.txt').T
6 k = 4 # dimension of latent

```

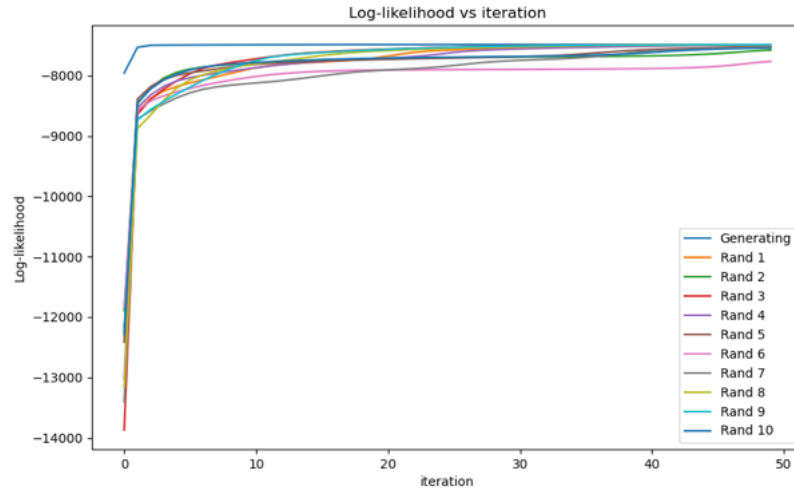


Figure 25: Log-likelihood vs iteration for generating and random initialisations over 50 iterations - and increase in likelihood and convergence is seen

```

7 d, T = X.shape #dimension data and number of time steps
8
9 #defining A
10 t1 = 2 * np.pi / 180
11 t2 = 2 * np.pi / 90
12 A = 0.99 * np.array([
13     [np.cos(t1), -np.sin(t1), 0, 0],
14     [np.sin(t1), np.cos(t1), 0, 0],
15     [0, 0, np.cos(t2), -np.sin(t2)],
16     [0, 0, np.sin(t2), np.cos(t2)]
17 ])
18 #Q
19 Q = np.eye(k) - A @ A.T
20 #C
21 C = np.array([
22     [1, 0, 1, 0],
23     [0, 1, 0, 1],
24     [1, 0, 0, 1],
25     [0, 0, 1, 1],
26     [0.5, 0.5, 0.5, 0.5]])
27 R = np.eye(d)
28
29 #initial states taken as mean of normal - 0 and Identity -
was not sure if I should instead sample from the normal
dist ect. but ended up converging to similar results
30 y_init = np.zeros(k)

```



```

31 Q_init = np.eye(k)
32
33 def run_ssm_kalman(X, y_init, Q_init, A, Q, C, R, mode='
smooth'):
34     """
35     Calculates kalman-smoother estimates of SSM state
        posterior.
36     :param X:      data, [d, t_max] numpy array
37     :param y_init:  initial latent state, [k,] numpy array
38     :param Q_init:  initial variance, [k, k] numpy array
39     :param A:      latent dynamics matrix, [k, k] numpy
        array
40     :param Q:      innovations covariance matrix, [k, k]
        numpy array
41     :param C:      output loading matrix, [d, k] numpy
        array
42     :param R:      output noise matrix, [d, d] numpy array
43     :param mode:    'forw' or 'filt' for forward filtering,
        'smooth' for also backward filtering
44     :return:
45     y_hat:         posterior mean estimates, [k, t_max] numpy
        array
46     V_hat:         posterior variances on y_t, [t_max, k, k]
        numpy array
47     V_joint:       posterior covariances between y_{t+1}, y_t,
        [t_max, k, k] numpy array
48     likelihood:    conditional log-likelihoods log(p(x_t|x_{1:t
        -1})), [t_max,] numpy array
49     """
50     d, k = C.shape
51     t_max = X.shape[1]
52
53     # dimension checks
54     assert np.all(X.shape == (d, t_max)), "Shape of X must
        be (%d, %d), %s provided" % (d, t_max, X.shape)
55     assert np.all(y_init.shape == (k,)), "Shape of y_init
        must be (%d,), %s provided" % (k, y_init.shape)
56     assert np.all(Q_init.shape == (k, k)), "Shape of Q_init
        must be (%d, %d), %s provided" % (k, k, Q_init.shape)
57     assert np.all(A.shape == (k, k)), "Shape of A must be (%
        d, %d), %s provided" % (k, k, A.shape)
58     assert np.all(Q.shape == (k, k)), "Shape of Q must be (%
        d, %d), %s provided" % (k, k, Q.shape)
59     assert np.all(C.shape == (d, k)), "Shape of C must be (%
        d, %d), %s provided" % (d, k, C.shape)
60     assert np.all(R.shape == (d, d)), "Shape of R must be (%
        d, %d), %s provided" % (d, k, R.shape)
61
62     y_filt = np.zeros((k, t_max)) # filtering estimate: \
        hat(y)_t^t

```

```

63 V_filt = np.zeros((t_max, k, k)) # filtering variance:
    \hat(V)_t^t
64 y_hat = np.zeros((k, t_max)) # smoothing estimate: \hat
    (y)_t^T
65 V_hat = np.zeros((t_max, k, k)) # smoothing variance: \
    hat(V)_t^T
66 K = np.zeros((t_max, k, X.shape[0])) # Kalman gain
67 J = np.zeros((t_max, k, k)) # smoothing gain
68 likelihood = np.zeros(t_max) # conditional log-
    likelihood: p(x_t/x_{1:t-1})
69
70 I_k = np.eye(k)
71
72 # forward pass
73
74 V_pred = Q_init
75 y_pred = y_init
76
77 for t in range(t_max):
78     x_pred_err = X[:, t] - C.dot(y_pred)
79     V_x_pred = C.dot(V_pred.dot(C.T)) + R
80     V_x_pred_inv = np.linalg.inv(V_x_pred)
81     likelihood[t] = -0.5 * (np.linalg.slogdet(2 * np.pi
        * (V_x_pred))[1] +
82                             x_pred_err.T.dot(
                                V_x_pred_inv).dot(
                                    x_pred_err))
83
84     K[t] = V_pred.dot(C.T).dot(V_x_pred_inv)
85
86     y_filt[:, t] = y_pred + K[t].dot(x_pred_err)
87     V_filt[t] = V_pred - K[t].dot(C).dot(V_pred)
88
89     # symmetrise the variance to avoid numerical drift
90     V_filt[t] = (V_filt[t] + V_filt[t].T) / 2.0
91
92     y_pred = A.dot(y_filt[:, t])
93     V_pred = A.dot(V_filt[t]).dot(A.T) + Q
94
95 # backward pass
96
97 if mode == 'filt' or mode == 'forw':
98     # skip if filtering/forward pass only
99     y_hat = y_filt
100     V_hat = V_filt
101     V_joint = None
102 else:
103     V_joint = np.zeros_like(V_filt)
104     y_hat[:, -1] = y_filt[:, -1]
105     V_hat[-1] = V_filt[-1]

```

```

106
107     for t in range(t_max - 2, -1, -1):
108         J[t] = V_filt[t].dot(A.T).dot(np.linalg.inv(A.
109             dot(V_filt[t]).dot(A.T) + Q))
110         y_hat[:, t] = y_filt[:, t] + J[t].dot((y_hat[:,
111             t + 1] - A.dot(y_filt[:, t])))
112         V_hat[t] = V_filt[t] + J[t].dot(V_hat[t + 1] - A
113             .dot(V_filt[t]).dot(A.T) - Q).dot(J[t].T)
114
115     V_joint[-2] = (I_k - K[-1].dot(C)).dot(A).dot(V_filt
116         [-2])
117
118     for t in range(t_max - 3, -1, -1):
119         V_joint[t] = V_filt[t + 1].dot(J[t].T) + J[t +
120             1].dot(V_joint[t + 1] - A.dot(V_filt[t + 1]))
121         .dot(J[t].T)
122
123     return y_hat, V_hat, V_joint, likelihood
124
125 #running kalman for filt
126 y_filt, V_filt, _, L_filt = run_ssm_kalman(X, y_init, Q_init
127     , A, Q, C, R, mode='filt')
128
129 #plotting filtered means
130 plt.figure(figsize=(10, 6))
131 for i in range(k):
132     plt.plot(y_filt[i, :], label=f'Latent {i+1}')
133 plt.title('filtered latent')
134 plt.xlabel('time')
135 plt.ylabel('value')
136 plt.legend()
137 plt.show()
138
139 #running kalman for smooth
140 y_smooth, V_smooth, _, L_smooth = run_ssm_kalman(X, y_init,
141     Q_init, A, Q, C, R, mode='smooth')
142
143 # plotting smooth mean
144 plt.figure(figsize=(10, 6))
145 for i in range(k):
146     plt.plot(y_smooth[i, :], label=f'Latent {i+1}')
147 plt.title('smoothed latent')
148 plt.xlabel('time')
149 plt.ylabel('value')
150 plt.legend()
151 plt.show()
152
153 # plot the log-determinants
154 plt.figure(figsize=(10, 6))
155 plt.plot(np.linalg.slogdet(V_filt)[1])
156 plt.title('Log-Determinant filt')

```

```

148 plt.xlabel('time')
149 plt.ylabel('ld')
150 plt.show()
151
152 plt.figure(figsize=(10, 6))
153 plt.plot(np.linalg.slogdet(V_smooth)[1])
154 plt.title('Log-Determinant smooth')
155 plt.xlabel('Time')
156 plt.ylabel('ld')
157 plt.show()
158
159 def em_lgssm(X, y_init, Q_init, A_init, Q_init_param, C_init
160             , R_init, num_iter=50):
161     """
162     Run EM algorithm to estimate parameters of LGSSM using E
163     and M updates pre calculated
164     params
165     X: data [d,t]
166     y_init: initial latent state
167     Q_init: initial variance
168     A_init: initial latent dynamics matrix [k, k]
169     Q_init_param: initial innovations covariance matrix [
170     k, k]
171     C_init: initial output loading matrix [d, k]
172     R_init: output noise matrix [d, d]
173     num_iter: number of iterations
174     output
175     A: updated A param
176     Q: updated q param
177     C: updated C param
178     R: updated R param
179     llh: log likelihoods stored each iteration
180     """
181     # Initialize parameters
182     A = A_init
183     Q = Q_init_param
184     C = C_init
185     R = R_init
186
187     # Store log-likelihoods
188     llh = []
189
190     for iteration in range(num_iter):
191         # E - running kalman smoother and log likelihoods
192         y_hat, V_hat, V_joint, likelihoods = run_ssm_kalman(
193             X, y_init, Q_init, A, Q, C, R, mode='smooth')
194
195         total_llh = np.sum(likelihoods)
196         llh.append(total_llh)

```

```

194     #M
195     T = X.shape[1]
196     k = y_hat.shape[0]
197     d = X.shape[0]
198
199     xx = np.zeros((d, d))
200     xy = np.zeros((d, k))
201     yy = np.zeros((k, k))
202     y_y_prev = np.zeros((k, k))
203     yy_prev = np.zeros((k, k))
204     yy_t = np.zeros((k, k))
205
206     for t in range(T):
207         x_t = X[:, t][:, np.newaxis] #d
208         y_t = y_hat[:, t][:, np.newaxis] #k
209         V_t = V_hat[t] #k x k
210
211         xx += x_t @ x_t.T
212         xy += x_t @ y_t.T
213
214         yy += y_t @ y_t.T + V_t
215
216     for t in range(1, T):
217         y_t = y_hat[:, t][:, np.newaxis] # k
218         y_prev = y_hat[:, t-1][:, np.newaxis] # k
219         V_t = V_hat[t] # k x k
220         V_prev = V_hat[t-1] # k x k
221         V_joint_t = V_joint[t-1] # k x k
222
223         y_y_prev += y_t @ y_prev.T + V_joint_t
224         yy_prev += y_prev @ y_prev.T + V_prev
225         yy_t += y_t @ y_t.T + V_t
226
227         # Update all parameters as in M step
228         C = xy @ np.linalg.inv(yy)
229         A = y_y_prev @ np.linalg.inv(yy_prev)
230         R = (1/T) * (xx - xy @ C.T)
231         Q = (1/(T-1)) * (yy_t - y_y_prev @ A.T)
232     return A, Q, C, R, llh
233     #use generating parameters
234     A_gen, Q_gen, C_gen, R_gen, llh_gen = em_lgssm(X, y_init,
235         Q_init, A, Q, C, R, num_iter=50)
236
237     #random initialisations - i randomised the A and C matrices
238     and let Q and R be identities - was unsure if I should
239     randomise everything?
240     llh_random_runs = []
241
242     for i in range(10):
243         A_rand = np.random.randn(k, k)

```

```

241 Q_rand = np.eye(k)
242 C_rand = np.random.randn(d, k)
243 R_rand = np.eye(d)
244
245 A_rand_n, Q_rand_n, C_rand_n, R_rand_n, llh_rand =
    em_lgssm(X, y_init, Q_init, A_rand, Q_rand, C_rand,
    R_rand, num_iter=50)
246
247 llh_random_runs.append(llh_rand)
248
249 plt.figure(figsize=(10, 6))
250
251 # plotting log likelihoods
252 plt.plot(llh_gen, label='Generating')
253 for idx, llh_rand in enumerate(llh_random_runs):
254     plt.plot(llh_rand, label=f'Rand {idx+1}')
255
256 plt.title('Log-likelihood vs iteration')
257 plt.xlabel('iteration')
258 plt.ylabel('Log-likelihood')
259 plt.legend()
260 plt.show()

```

Listing 4: LGSSM python code and plots

5 Decrypting Messages with MCMC

a) We model the English text $s_1 s_2 \dots$ as a first order Markov chain such that

$$p(s_1 s_2 \dots s_n) = p(s_1) \prod_{i=2}^n p(s_i | s_{i-1})$$

We let the stationary distribution be represented as $\phi(\gamma) = \lim_{i \rightarrow \infty} p(s_i = \gamma)$ and transition probabilities as $p(s_i = \alpha | s_{i-1} = \beta) = \psi(\alpha, \beta)$.

The stationary distribution $\phi(\gamma)$ represents the long-run proportion of times symbol γ appears in the text. In the limit where you have a very large dataset the probability tends to:

$$\phi(\gamma) = \frac{N_\gamma}{N}$$

where N is the total number of symbols in the text:

$$N = \sum_{\gamma} N_\gamma$$

The ML estimate for $\psi(\alpha, \beta)$ is given by:




Figure 27: Phi probabilities - CDF

Hence, for $s_i = \sigma^{-1}(e_1)$ i.e. the decryption

$$p(e_1 e_2 \dots | \sigma) = \phi(\sigma^{-1}(e_1)) \prod_{i=2}^n \psi(\sigma^{-1}(e_i) | \sigma^{-1}(e_{i-1}))$$

c) We define a function $S(\sigma \rightarrow \sigma')$ that represents a swap of 2 of the symbols at random from 2 permutations σ and σ' . There is $53C2$ different pairings of symbols among the total of 53 where C represents a combination. Given that we are restricted by this action we have that $S(\sigma \rightarrow \sigma') = \frac{1}{53C2}$ for $\sigma'(s_i) = \sigma(s_j)$ and $\sigma'(s_j) = \sigma(s_i)$ and otherwise $S(\sigma \rightarrow \sigma') = 0$.

The acceptance in lectures is defined as

$$A = \min\left\{1, \frac{S(\sigma' \rightarrow \sigma)P(\sigma' | e_1 \dots e_n)}{S(\sigma \rightarrow \sigma')P(\sigma | e_1 \dots e_n)}\right\}$$

However we know that S is a symmetric distribution so these cancel out. Using Bayes theorem we also have that

$$P(\sigma' | e_1 \dots e_n) = \frac{P(e_1 \dots e_n | \sigma')P(\sigma')}{\sum_i P(e_1 \dots e_n | \sigma_i)P(\sigma_i)}$$

$P(\sigma | e_1 \dots e_n)$ produces something very similar and given the denominator is a constant and we have a uniform prior, $P(\sigma) = P(\sigma')$, this means we can reduce the acceptance to:

$$A = \min\left\{1, \frac{P(e_1 \dots e_n | \sigma')}{P(e_1 \dots e_n | \sigma)}\right\}$$

d) The proposal distributions were sampled from the probabilities calculated from the war and peace text as this seemed to be a large dataset in comparison

to our encrypted text and one that was reasonable to use as a proposal. Initially, I had random initialisation which sort of worked giving a somewhat accurate answer after approximately 8-10 tries.

I then tried initialising such that each of the characters in the ϕ table matched to the encrypted text from most common to least common in the encrypted matching up with the highest to lowest probabilities from the large text. This however produced some non-sensical results as I think the model was falling into an incorrect local stationary point or because there were symbols that did not appear in the encrypted text or had 0 probabilities for ϕ .

In order to avoid this I added 1 to all of the frequencies meaning that no symbol had an exactly 0 probability and I mapped the space to the most common symbol and then after testing a few different combinations and varying the number of symbols mapped I picked the next 5 to the top 5 most common symbols. This allowed for some initialisation as well as randomness and seemed to work much better - getting a reasonable answer around every 4 runs.

```
1 import numpy as np
2 import random
3 import pandas as pd
4 from matplotlib import pyplot as plt
5 # reading the file
6 with open('/Users/baidn/Downloads/symbols.txt', 'r') as f:
7     symbols = [line.rstrip('\n') for line in f]
8     symbols = [' ' if symbol == '' else symbol for symbol in
9                symbols]
10    #print("Symbols:", symbols)
11
12    # symbol to index and index to symbol mappings
13    symbol_to_i = {s: i for i, s in enumerate(symbols)}
14    i_to_symbol = {i: s for i, s in enumerate(symbols)}
15    K = len(symbols)
16
17    # read war and peace for calculating proposals
18    with open('/Users/baidn/Downloads/WarAndPeace.txt', 'r',
19              encoding='utf-8') as f:
20        text = f.read()
21    #ignoring the intro text - only considering the book
22    text = text[834:].lower()
23    #count symbols and pairs in loop and calc probability
24    symbol_c = np.zeros(K, dtype=np.int64)
25    pair_c = np.zeros((K, K), dtype=np.int64)
26    prev_idx = None
27    for s in text:
28        if s in symbol_to_i:
29            idx = symbol_to_i[s]
30            symbol_c[idx] += 1
31            if prev_idx is not None:
32                pair_c[prev_idx, idx] += 1
33            prev_idx = idx
```

```

32     else:
33         continue # skip symbols not in our set
34 symbol_c[symbol_c == 0] = 1 #make probabilities slightly
    more than 0 for symb that dont occur such as not to leave
    these out - leads to better results
35 N = symbol_c.sum()
36 phi_arr = symbol_c / N
37
38 symbol_counts_no_zeros = np.where(symbol_c == 0, 1, symbol_c
    )
39 psi_arr = pair_c / symbol_counts_no_zeros[:, None]
40 psi_arr = np.nan_to_num(psi_arr)
41 print("Psi array shape:", psi_arr.shape)
42 print("Psi array:", psi_arr)
43 print(phi_arr)
44 #heatmap for the transition matrix and a graph for phi
45 char = ['space' if s == ' ' else s for s in symbols]
46
47 plt.figure(figsize=(12, 10))
48 plt.imshow(psi_arr, cmap='viridis', interpolation='nearest')
49 plt.colorbar()
50 plt.xticks(ticks=np.arange(K), labels=char, rotation=90,
    fontsize=8)
51 plt.yticks(ticks=np.arange(K), labels=char, fontsize=8)
52 plt.title('Transition Matrix')
53 plt.xlabel('Next symbol')
54 plt.ylabel('Current symbol')
55 plt.tight_layout()
56 plt.show()
57
58 plt.figure(figsize=(12, 6))
59 plt.bar(char, phi_arr)
60 plt.xlabel('Symb')
61 plt.ylabel('Phi Prob')
62
63 plt.xticks(rotation=90)
64 plt.tight_layout()
65 plt.show()
66
67 epsilon = 1e-10 #to prevent log(0)
68 #log probabilities
69 log_phi = np.log(phi_arr + epsilon)
70 log_psi = np.log(psi_arr + epsilon)
71
72 # read encrypted text
73 with open('/Users/baidn/Downloads/message.txt', 'r') as f:
74     encrypt_txt = f.read().strip()
75
76 # encrypted symbols to indices
77 encrypt_i = [symbol_to_i[s] for s in encrypt_txt if s in

```

```

    symbol_to_i]
78 sigma = -np.ones(K, dtype=int) # initializing sigma
79 # keep track of assigned characters
80 cipher_i = set()
81 plaintext_i = set()
82
83 com_char = [' ', 'e', 'a', 't', 'i', 'n']
84
85 # get indices of common characters
86 com_char_i = [symbol_to_i[c] for c in com_char if c in
    symbol_to_i]
87
88 # Sort encrypted symbols by frequency in descending order
89 sort_encrypt_i = sorted(range(K), key=lambda x: -symbol_c[x
    ])
90
91 # Select the top frequent encrypted symbols and map to
    common characters
92 top_encrypt = sort_encrypt_i[:len(com_char_i)]
93 for enc_idx, plain_idx in zip(top_encrypt, com_char_i):
94     sigma[enc_idx] = plain_idx
95     cipher_i.add(enc_idx)
96     plaintext_i.add(plain_idx)
97
98 # remaining cipher and plaintext indices and random mapping
    of these
99 r_cipher_i = set(range(K)) - cipher_i
100 r_plaintext_i = set(range(K)) - plaintext_i
101 r_cipher_i = list(r_cipher_i)
102 r_plaintext_i = list(r_plaintext_i)
103 random.shuffle(r_plaintext_i)
104
105 for enc_idx, avail_idx in zip(r_cipher_i, r_plaintext_i):
106     sigma[enc_idx] = avail_idx
107     cipher_i.add(enc_idx)
108     plaintext_i.add(avail_idx)
109
110 #inverse
111 sigma_inv = np.zeros(K, dtype=int)
112 sigma_inv[sigma.astype(int)] = np.arange(K)
113
114 #log likelihood
115 def compute_log_likelihood(decrypt_i):
116     log_ml = log_phi[decrypt_i[0]]
117     for i in range(1, len(decrypt_i)):
118         prev_idx = decrypt_i[i - 1]
119         curr_idx = decrypt_i[i]
120         log_ml += log_psi[prev_idx, curr_idx]
121     return log_ml
122

```

```

123 decrypt_i = [sigma_inv[idx] for idx in encrypt_i]
124 log_p_e_g_sigma = compute_log_likelihood(decrypt_i)
125
126
127 # MH sampler and tracking acceptance rate
128 iter = 10000
129 acceptance = 0
130 for iteration in range(1, iter + 1):
131     #new sigma by swapping two symbols
132     i, j = random.sample(range(K), 2)
133     sigma_prop = sigma.copy()
134     sigma_prop[i], sigma_prop[j] = sigma_prop[j], sigma_prop
        [i]
135
136     #update sigma_inverse accordingly
137     sigma_i_prop = np.zeros(K, dtype=int)
138     sigma_i_prop[sigma_prop.astype(int)] = np.arange(K)
139
140     #decrypt the message with proposed sigma
141     decrypted_i_prop = [sigma_i_prop[idx] for idx in
        encrypt_i]
142
143     #new log-likelihood, use logs to prevent some small
        values being treated as zeros
144     log_p_e_g_sigma_prop = compute_log_likelihood(
        decrypted_i_prop)
145
146     #acceptance probability
147     del_log_ml = log_p_e_g_sigma_prop - log_p_e_g_sigma
148     acceptance_prob = min(1, np.exp(del_log_ml))
149
150     # accpet/reject step
151     if random.random() < acceptance_prob:
152         sigma = sigma_prop
153         sigma_inv = sigma_i_prop
154         decrypt_i = decrypted_i_prop
155         log_p_e_g_sigma = log_p_e_g_sigma_prop
156         acceptance += 1
157
158
159     # 100 iterations print
160     if iteration % 100 == 0:
161         decrypted_text_sample = ''.join([i_to_symbol[idx]
            for idx in decrypt_i[:60]])
162         print(f"Iteration {iteration}: {
            decrypted_text_sample}")
163 print(f"overall Acceptance Rate: {acceptance / iter}")

```

Listing 5: MCMC Decryption python code and transition matrix and phi plots

```

/Users/baidu/Downloads/Decrypting_Messages_with_MONC.py:528: RuntimeWarning: o
acceptance_probability = min(1, np.exp(delta_log_likelihood))
Iteration 100: 7 .ua.ah" xis.l n.uhi.n"z isl/zl.allsb.ua.jleois.xlmi.ui.bh
Iteration 200: 7 .-b.bd" xis.o n.-dsic"z isozlibiosa.-b.koelis.xoci.-l.ad
Iteration 300: u l-bliber xislo nl-dsiffr isozlibiosa.-l-blooisixofil.-ll.d
Iteration 400: u l-blibd. xirto nl-drilfr.s irozlibiorwl.-blooisirxofil.-llwd
Iteration 500: u l-blibe. garlo nl-eralf.m arozmalbaerwl.-bkodiarlgofal.-alwe
Iteration 600: u l-blibo. garle nl-oralk.m arezmalbaerwl.-blfediargheal.-alwo
Iteration 700: u l-blibok garle nl-oralk.m arezmalbaerwl.-blfediargheal.-alwo
Iteration 800: s l-blibok garle nl-oralk.m arezmalbaerwl.-blfediargheal.-alwo
Iteration 900: e l-blibok garls nl-oralk.m arszmalbasrwl.-blfstiarlgs.al.-alwo
Iteration 1000: e l-blibok garls wl-oralk.m arszmalbasrwl.-blfstiarlgs.al.-alno
Iteration 1100: e lyblbok garls wlyoral.km arszmalbasrwlblfstiarlgs.alyalno
Iteration 1200: e lyblbok garls wlyoral.km arszmalbasrwlblfstiarlgs.alyalno
Iteration 1300: e lyblbou warls glyoralcum arszmalbasrwlblfstiarlgs.alyalno
Iteration 1400: e lyblbou warls glyoralcum arszmalbasrwlblfstiarlgs.alyalno
Iteration 1500: e lyblbou warls glyoralcum arszmalbasrwlblfstiarlgs.alyalno
Iteration 1600: e lyblbou warls glyoralcum arszmalbasrwlblfstiarlgs.alyalno
Iteration 1700: e lyblbou warls glyoralcum arszmalbasrwlblfstiarlgs.alyalno
Iteration 1800: e lyblbou warls glyoralcum arszmalbasrwlblfstiarlgs.alyalno
Iteration 1900: e lyblbau corls glyoralhum ors-molbosrnlblksitorlcscholyolna
Iteration 2000: e lyblbau corls glyoralhum ors-molbosrnlblksitorlcscholyolna
Iteration 2100: e lyblbau corls glyoralhum ors-molbosrnlblksitorlcscholyolna
Iteration 2200: e lyblbau corls glyoralhum ors-molbosrnlblksitorlcscholyolna
Iteration 2300: e lyblbau corls glyoralhum ors-molbosrnlblksitorlcscholyolna
Iteration 2400: e lyblbau corls glyoralhum ors-molbosrnlblksitorlcscholyolna
Iteration 2500: e lyblbau corls glyoralhum ors-molbosrnlblksitorlcscholyolna
Iteration 2600: e lyblbau corls glyoralhum ors-molbosrnlblksitorlcscholyolna
Iteration 2700: el y. .aulcor ilf yaro humlorismo .oirn y. kistor cibe yo na
Iteration 2800: ol y. .aulcor ilf yare bunleriyne .eirn y. kister cibe ye na
Iteration 2900: ol y. .aulcor ilf yare bunleriyne .eirn y. kister cibe ye na
Iteration 3000: ol w. .aulcor ilf ware bunleriyne .eirn w. kister cibe we na
Iteration 3100: ol w. .aulcor ilf ware bunleriyne .eirn w. kister cibe we na
Iteration 3200: ol w. .aulcor ilf ware bunleriyne .eirn w. kister cibe we na
Iteration 3300: al w. .doulker ilf wore bunleriyne .eirn w. gisher kibe we no
Iteration 3400: al wd .doulker ilf wore bunleriyne .eirn wd gisher kibe we no
Iteration 3500: al wd .doulker ilf wore bunleriyne .eirn wd gisher kibe we no
Iteration 3600: al tn .noulker ilf tore wutlerikne .neird tn gisher yive te do
Iteration 3700: al mn .noulker ilf more wutlerikne .neird mn gisher yive me do
Iteration 3800: al mn .noulker ilg more wutlerikne .neird mn fisher yive me do
Iteration 3900: al mn .noulker ilg more wutlerikne .neird mn fisher yive me do
Iteration 4000: al mn .noulker ilg more wutlerikne .neird mn fisher yive me do
Iteration 4100: al mn .noulker ilg more wutlerikne .neird mn fisher yive me do
Iteration 4200: al wy .younler ilg wore mulderibde yeirs wy fither nine we so
Iteration 4300: il wy .younler alg wore mulderibde years wy father name we so
Iteration 4400: il wy .younler alg wore mulderibde years wy father name we so
Iteration 4500: id wy .younder adg wore mulderable years wy father name we so

```

Figure 28: First 60 decrypted characters -part 1

```

Iteration 4600: id wy .younder adg wore mulderable years wy father name we so
Iteration 4700: id wy .younder adg wore mulderable years wy father name we so
Iteration 4800: id wy .younder adg wore mulderable years wy father name we so
Iteration 4900: in my .younder ang more vulnerable years my father dawe me so
Iteration 5000: in my .younder ang more vulnerable years my father dawe me so
Iteration 5100: in my .younder ang more vulnerable years my father dawe me so
Iteration 5200: in my .younder ang more vulnerable years my father dawe me so
Iteration 5300: in my .younder ang more vulnerable years my father dawe me so
Iteration 5400: in my .younder ang more vulnerable years my father dawe me so
Iteration 5500: in my .younder ang more vulnerable years my father dawe me so
Iteration 5600: in my .younder ang more vulnerable years my father dawe me so
Iteration 5700: in my .younder ang more vulnerable years my father dawe me so
Iteration 5800: in my .younder ang more vulnerable years my father dawe me so
Iteration 5900: in my .younder ang more vulnerable years my father dawe me so
Iteration 6000: in my .younder ang more vulnerable years my father dawe me so
Iteration 6100: in my .younder ang more vulnerable years my father dawe me so
Iteration 6200: in my .younder ang more vulnerable years my father dawe me so
Iteration 6300: in my .younder ang more vulnerable years my father dawe me so
Iteration 6400: in my .younder ang more vulnerable years my father dawe me so
Iteration 6500: in my .younder ang more vulnerable years my father dawe me so
Iteration 6600: in my .younder ang more vulnerable years my father dawe me so
Iteration 6700: in my .younder ang more vulnerable years my father dawe me so
Iteration 6800: in my .younder ang more vulnerable years my father dawe me so
Iteration 6900: in my .younder ang more vulnerable years my father dawe me so
Iteration 7000: in my .younder ang more vulnerable years my father dawe me so
Iteration 7100: in my .younder ang more vulnerable years my father dawe me so
Iteration 7200: in my .younder ang more vulnerable years my father dawe me so
Iteration 7300: in my .younder ang more vulnerable years my father dawe me so
Iteration 7400: in my .younder ang more vulnerable years my father dawe me so
Iteration 7500: in my .younder ang more vulnerable years my father dawe me so
Iteration 7600: in my .younder ang more vulnerable years my father dawe me so
Iteration 7700: in my .younder ang more vulnerable years my father dawe me so
Iteration 7800: in my .younder ang more vulnerable years my father dawe me so
Iteration 7900: in my .younder ang more vulnerable years my father dawe me so
Iteration 8000: in my .younder ang more vulnerable years my father dawe me so
Iteration 8100: in my .younder ang more vulnerable years my father dawe me so
Iteration 8200: in my .younder ang more vulnerable years my father dawe me so
Iteration 8300: in my .younder ang more vulnerable years my father dawe me so
Iteration 8400: in my .younder ang more vulnerable years my father dawe me so
Iteration 8500: in my .younder ang more vulnerable years my father dawe me so
Iteration 8600: in my .younder ang more vulnerable years my father dawe me so
Iteration 8700: in my .younder ang more vulnerable years my father dawe me so
Iteration 8800: in my .younder ang more vulnerable years my father dawe me so
Iteration 8900: in my .younder ang more vulnerable years my father dawe me so
Iteration 9000: in my .younder ang more vulnerable years my father dawe me so

```

Figure 29: First 60 decrypted characters -part 2

e) The initial set up is not ergodic as this would imply you could reach any state from any other state through some path and the chain is aperiodic. Due to the zero values I would have thought there is still potential to reach x from y even if $\psi(x, y) = 0$ through a series of swaps. However, I think there are 0 columns in the transition matrix which means we will not be able to reach certain states as there is a 0 chance that state x will go to y for example if $\psi(z, y) = 0$ for z being all characters. We can fix this by adding small probabilities to all of the 0 values such that the probability is never 0 i.e. start with $N_\gamma = 1$.

f) Symbol probabilities are not sufficient as it would converge to mapping the most probable symbols in each of the others. This would remove all contextual information which is crucial to language. Also, there would be ambiguity if frequency of symbols matched.

Using a second-order Markov chain introduces challenges, including data sparsity, as there is an increase in zero probabilities, and computational complexity due to the number of transition probabilities growing to S^3 (for S the number of symbols) which would also require a larger training dataset. Additionally, ergodicity concerns arise due to more zero transitions, contributing to irreducibility issues.

This approach fails if two symbols map to the same encrypted value as the model relies on a bijective mapping. This assumption of permutation fails when the mapping is not unique, leading to ambiguity in recovering the original symbols. With this approach we would have to have some probability associated with the possible mappings and it would be much harder to create a proposal.

The approach is impractical for languages with large symbol sets, such as Chinese, due to computational infeasibility and the large state space from over 10,000 symbols. The memory required for storing bigram probabilities is unmanageable, and there is insufficient data to accurately estimate frequencies and transitions. Also, the sampler could take a very long time to converge as it needs to explore this space.

6

7 Optimisation

a) To find the local extrema of the function $f(x, y) = x + 2y$ subject to the constraint $g(x, y) = y^2 + xy - 1 = 0$, we use the method of Lagrange multipliers.

Let λ be the Lagrange multiplier, and set up the Lagrangian function:

$$L(x, y, \lambda) = x + 2y - \lambda(y^2 + xy - 1)$$

The partial derivative with respect to x is

$$\frac{\partial L}{\partial x} = 1 - \lambda y = 0, \quad \Rightarrow \quad \lambda y = 1$$

With respect to y ,

$$\frac{\partial L}{\partial y} = 2 - \lambda(2y + x) = 0$$

The constraint equation is

$$g(x, y) = y^2 + xy - 1 = 0$$

b) From $\lambda = \frac{1}{y}$, we substitute into the second equation and solve for x and y . The local extrema occur at $(0, 1)$ and $(0, -1)$, where the function $f(x, y)$ attains its local extrema under the given constraint.

To find $x = \ln(a)$ using Newton's method, we consider the equation $\exp(x) = a$, rewritten as a root-finding problem:

$$f(x, a) = \exp(x) - a = 0$$

Newton's method updates x_n using:

$$x_{n+1} = x_n - \frac{\exp(x_n) - a}{\exp(x_n)} = x_n - 1 + \frac{a}{\exp(x_n)}$$

8 Eigenvalues as solutions of an optimisation problem

a) In this question $q_A = x^T A x$ and $R_A = \frac{x^T A x}{x^T x} = \frac{q_A}{|x|^2}$ and we aim to show that $\sup_{x \in \mathbb{R}^n} R_A(x)$ is attained. consider the unit sphere as \mathbb{R}^n is not contained and every \mathbb{R}^n can be converted to the unit sphere by $\frac{1}{|x|} = c$ for scalar c This is formally defined as:

$$S = \{x \in \mathbb{R}^n | |x| = 1\},$$

Note that for any nonzero scalar c ,

$$R_A(cx) = \frac{(cx)^T A (cx)}{(cx)^T (cx)} = \frac{c^2 x^T A x}{c^2 x^T x} = \frac{x^T A x}{x^T x} = R_A(x).$$

Noting that c is scalar its transpose is itself so $R_A(x)$ is homogeneous of degree zero and depends only on the direction of x .

Therefore,

$$\sup_{x \in \mathbb{R}^n} R_A(x) = \sup_{x \in S} R_A(x).$$

as $q_A(x)$ is continuous and $|x| = 1$ on S , $R_A(x)$ is continuous on S . By the Extreme Value Theorem, $R_A(x)$ attains its maximum on the compact set S .

b) A is symmetric with eigenvalues $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_n$ and corresponding orthonormal eigenvectors $\{\xi_1, \dots, \xi_n\}$.

$$x^T x = \sum_{i=1}^n (\xi_i^T x)^2,$$

Given that $A\xi_i = \lambda_i$

$$x^T Ax = \sum_{i=1}^n \lambda_i (\xi_i^T x)^2.$$

So the Rayleigh quotient becomes

$$R_A(x) = \frac{x^T Ax}{x^T x} = \frac{\sum_{i=1}^n \lambda_i (\xi_i^T x)^2}{\sum_{i=1}^n (\xi_i^T x)^2}.$$

Since $\lambda_i \leq \lambda_1$ for all i and given λ_1 is a constant we can take it out to give

$$\sum_{i=1}^n \lambda_i (\xi_i^T x)^2 \leq \sum_{i=1}^n \lambda_1 (\xi_i^T x)^2 = \lambda_1 \sum_{i=1}^n (\xi_i^T x)^2$$

dividing through

$$\begin{aligned} \frac{x^T Ax}{x^T x} &= \frac{\sum_{i=1}^n \lambda_i (\xi_i^T x)^2}{\sum_{i=1}^n (\xi_i^T x)^2} \leq \lambda_1 \\ \implies R_A(x) &\leq \lambda_1. \end{aligned}$$

(c) Suppose $x \in \mathbb{R}^n$ is not in $\text{span}\{\xi_1, \dots, \xi_k\}$, where $\lambda_1 = \dots = \lambda_k$ are the largest eigenvalues with multiplicity k .

Then there exists $j > k$ such that $\xi_j^T x \neq 0$. Using

$$R_A(x) = \frac{\sum_{i=1}^n \lambda_i (\xi_i^T x)^2}{\sum_{i=1}^n (\xi_i^T x)^2} = \lambda_1 - \frac{\sum_{i=k+1}^n (\lambda_1 - \lambda_i) (\xi_i^T x)^2}{\sum_{i=1}^n (\xi_i^T x)^2}.$$

Since $\lambda_i < \lambda_1$ for $i > k$ and $(\xi_i^T x)^2 > 0$ for some $i > k$, it follows that

$$\sum_{i=k+1}^n (\lambda_1 - \lambda_i) (\xi_i^T x)^2 > 0,$$

and hence,

$$R_A(x) = \lambda_1 - \frac{\text{something +ve}}{\sum_{i=1}^n (\xi_i^T x)^2} < \lambda_1.$$

Therefore, $R_A(x) < \lambda_1$ whenever $x \notin \text{span}\{\xi_1, \dots, \xi_k\}$.